

Machine learning workflows

MODELING WITH TIDYMODELS IN R



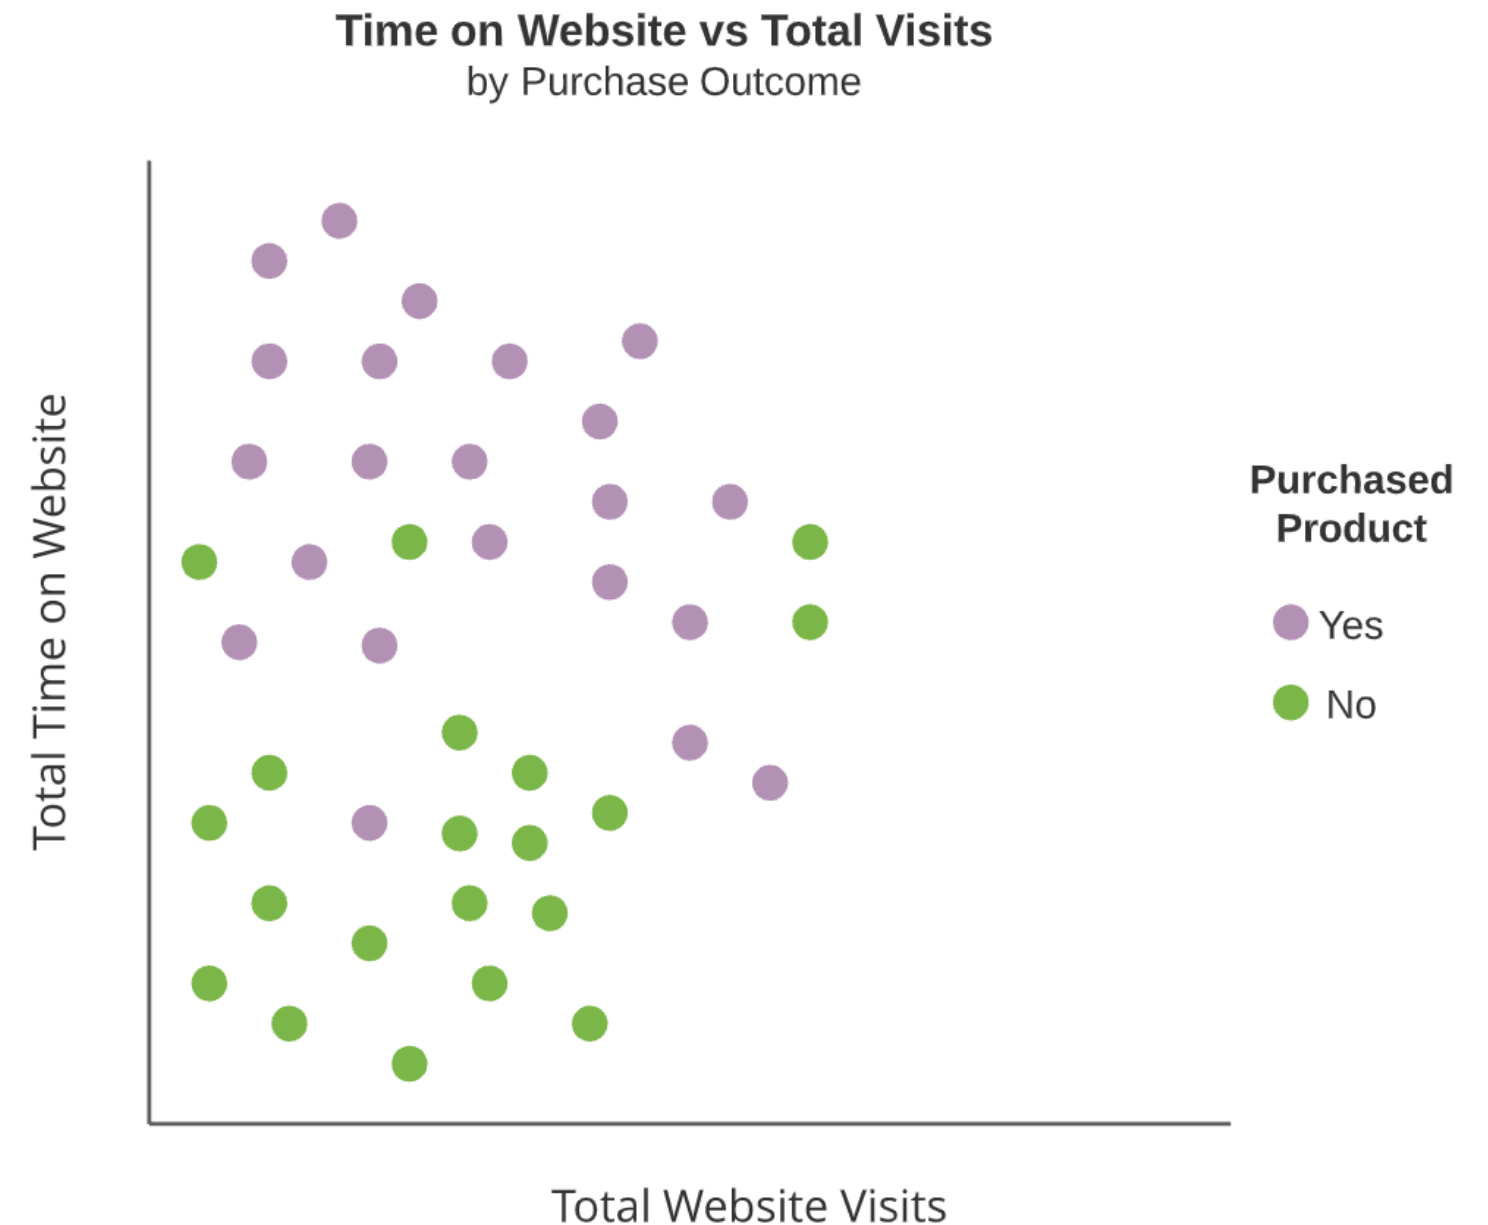
David Svancer
Data Scientist

Classification with decision trees

Decision trees segment the predictor space into **rectangular** regions

Recursive binary splitting

- Algorithm that segments predictor space into non-overlapping rectangular regions

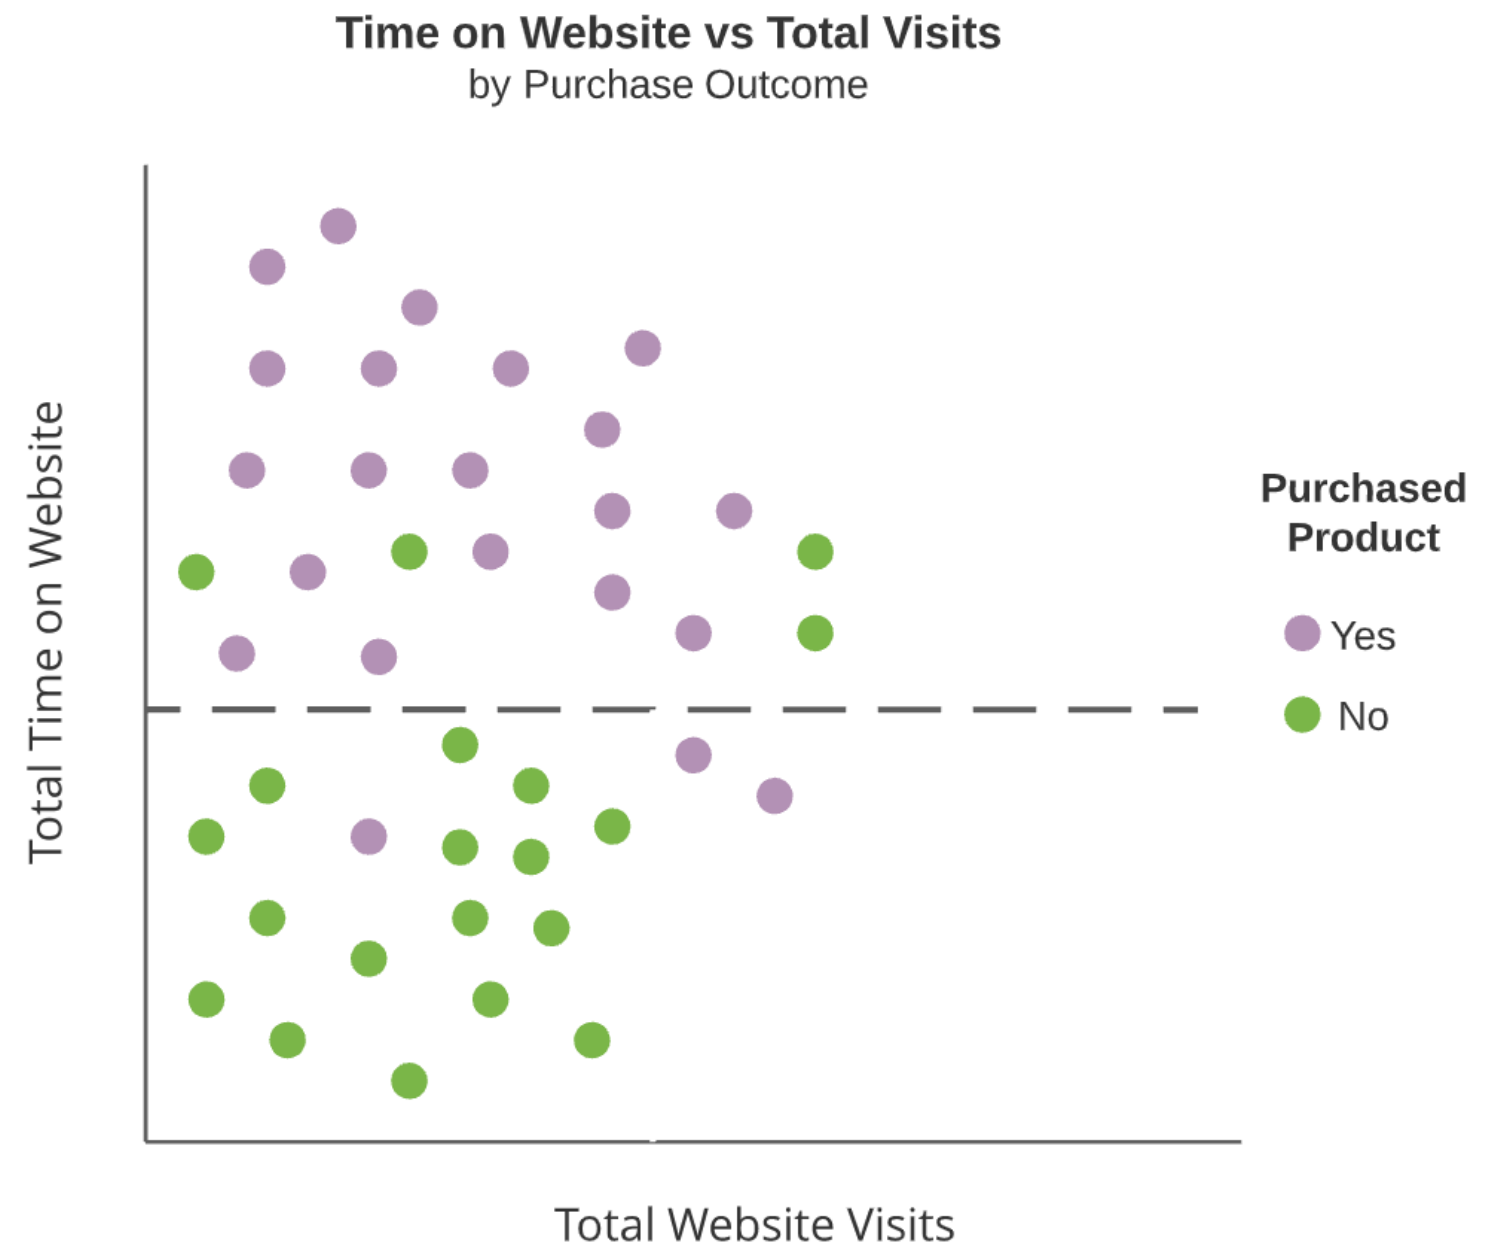


Classification with decision trees

Decision trees segment the predictor space into **rectangular** regions

Recursive binary splitting

- Algorithm that segments predictor space into non-overlapping rectangular regions
- Decision splits are added iteratively
 - Either horizontal or vertical cut points




Classification with decision trees

Decision trees segment the predictor space into **rectangular** regions

Recursive binary splitting

- Algorithm that segments predictor space into non-overlapping rectangular regions
- Decision splits are added iteratively
 - Either horizontal or vertical cut points




Classification with decision trees

Decision trees segment the predictor space into **rectangular** regions

Recursive binary splitting

- Algorithm that segments predictor space into non-overlapping rectangular regions
- Decision splits are added iteratively
 - Either horizontal or vertical cut points



Classification with decision trees


Decision trees segment the predictor space into **rectangular** regions

Recursive binary splitting

- Algorithm that segments predictor space into non-overlapping rectangular regions
- Decision splits are added iteratively
 - Either horizontal or vertical cut points

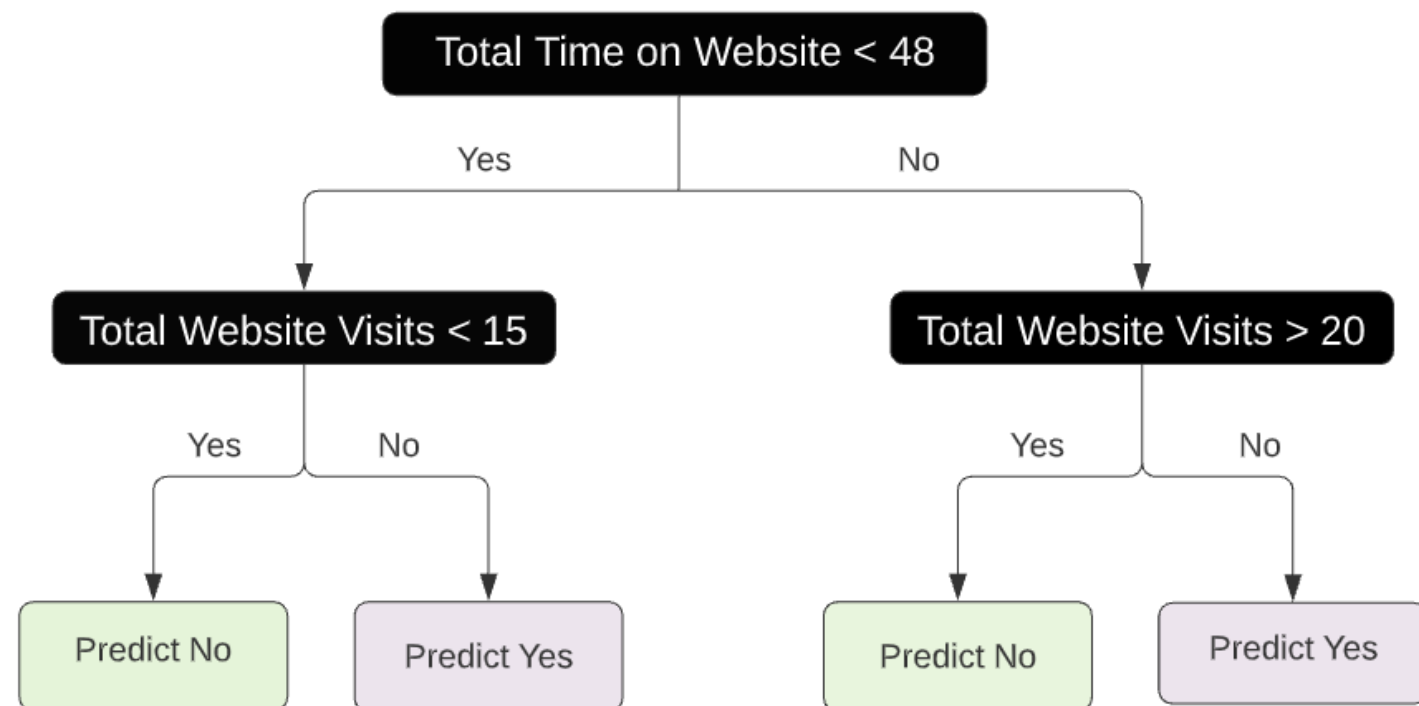
Produces distinct rectangular regions

- For classification, majority class is




Tree diagrams

- **Interior nodes**
 - Decision tree splits (dark boxes)
- **Terminal nodes**
 - Regions which are not split further
 - Green and purple boxes



Interior nodes are dashed lines and terminal nodes are highlighted rectangular regions



Model specification

Model specification in `parsnip`

- `decision_tree()`
 - General interface to decision tree models in `parsnip`
 - Common engine is 'rpart'
 - Mode can be either 'classification' or 'regression'
 - For lead scoring data, we need 'classification'

```
dt_model <- decision_tree() %>%  
  set_engine('rpart') %>%  
  set_mode('classification')
```


Feature engineering recipe

Data transformations for lead scoring data

- Encoded in a `recipe` object
 - Remove multicollinearity
 - Normalize numeric predictors
 - Create dummy variables for nominal predictors

Two R objects to manage

- `parsnip` model and `recipe` specification
- Combining into one object would make life easier

```
leads_recipe <- recipe(purchased ~ .,  
                        data = leads_training) %>%  
  step_corr(all_numeric(), threshold = 0.9) %>%  
  step_normalize(all_numeric()) %>%  
  step_dummy(all_nominal(), -all_outcomes())
```

```
leads_recipe
```

Data Recipe

Inputs:

	role	#variables
outcome		1
predictor		6

Operations:

Correlation filter on `all_numeric()`

Centering and scaling for `all_numeric()`

Dummy variables from `all_nominal()`, `-all_outcomes()`

Combining models and recipes

The `workflows` package is designed for streamlining the model process

- Combines a `parsnip` model and `recipe` object into a single `workflow` object

Initialized with the `workflow()` function

- Add model object with `add_model()`
- Add `recipe` object with `add_recipe()`
 - Must be specification, not a trained `recipe`

```
leads_wkfl <- workflow() %>%  
  add_model(dt_model) %>%  
  add_recipe(leads_recipe)
```

```
leads_wkfl
```

```
== Workflow =====  
Preprocessor: Recipe  
Model: decision_tree()  
-- Preprocessor -----  
3 Recipe Steps  
* step_corr()  
* step_normalize()  
* step_dummy()  
-- Model -----  
Decision Tree Model Specification (classification)  
Computational engine: rpart
```

Model fitting with workflows

Training a `workflow` object

- Pass `workflow` to `last_fit()` and provide data split object
- View model evaluation results with `collect_metrics()`

Behind the scenes

- Training and test datasets created
- `recipe` trained and applied
- Decision tree trained with training data
- Predictions and metrics on test data

```
leads_wkfl_fit <- leads_wkfl %>%  
  last_fit(split = leads_split)
```

```
leads_wkfl_fit %>%  
  collect_metrics()
```

```
# A tibble: 2 x 3  
  .metric .estimator .estimate  
  <chr>    <chr>         <dbl>  
1 accuracy binary        0.771  
2 roc_auc  binary        0.775
```

Collecting predictions

A `workflow` trained with `last_fit()` can be passed to `collect_predictions()`

- Produces detailed results on the test data
- Like before, can be used with `yardstick` functions to explore performance custom metrics

```
leads_wkfl_preds <- leads_wkfl_fit %>%  
  collect_predictions()
```

```
leads_wkfl_preds
```

```
# A tibble: 332 x 6  
  id          .pred_yes .pred_no  .row .pred_class purchased  
  <chr>         <dbl>   <dbl> <int>   <fct>      <fct>  
train/test split  0.120    0.880     2     no        no  
train/test split  0.755    0.245    17     yes       yes  
train/test split  0.120    0.880    21     no        no  
train/test split  0.120    0.880    22     no        no  
train/test split  0.755    0.245    24     yes       yes  
# ... with 327 more rows
```

Exploring custom metrics

Create a custom metric set with

`metric_set()`

- Area under the ROC curve, sensitivity, and specificity

Pass predictions datasets to

`leads_metrics()` to calculate metrics

```
leads_metrics <- metric_set(roc_auc, sens, spec)
```

```
leads_wkfl_preds %>%  
  leads_metrics(truth = purchased,  
                estimate = .pred_class,  
                .pred_yes)
```

```
# A tibble: 3 x 3  
  .metric .estimator .estimate  
  <chr>    <chr>         <dbl>  
1 sens    binary         0.75  
2 spec    binary         0.783  
3 roc_auc binary         0.775
```

Loan default dataset

Financial data for consumer loans at a bank

- Outcome variable is `loan_default`

```
loans_df
```

```
# A tibble: 872 x 8
  loan_default loan_purpose missed_payment_2_yr loan_amount interest_rate installment annual_income debt_to_income
  <fct>        <fct>        <fct>          <int>         <dbl>         <dbl>         <dbl>         <dbl>
1 no          debt_consolidation no           25000         5.47          855.         62823         39.4
2 yes         medical          no           10000        10.2          364.         40000         24.1
3 no          small_business no           13000         6.22          442.         65000         14.0
4 no          small_business no           36000         5.97         1152.        125000          8.09
5 yes         small_business yes           12000        11.8          308.         65000         20.1
# ... with 867 more rows
```

Let's practice building workflows!

MODELING WITH TIDYMODELS IN R

Estimating performance with cross validation

MODELING WITH TIDYMODELS IN R



David Svancer
Data Scientist

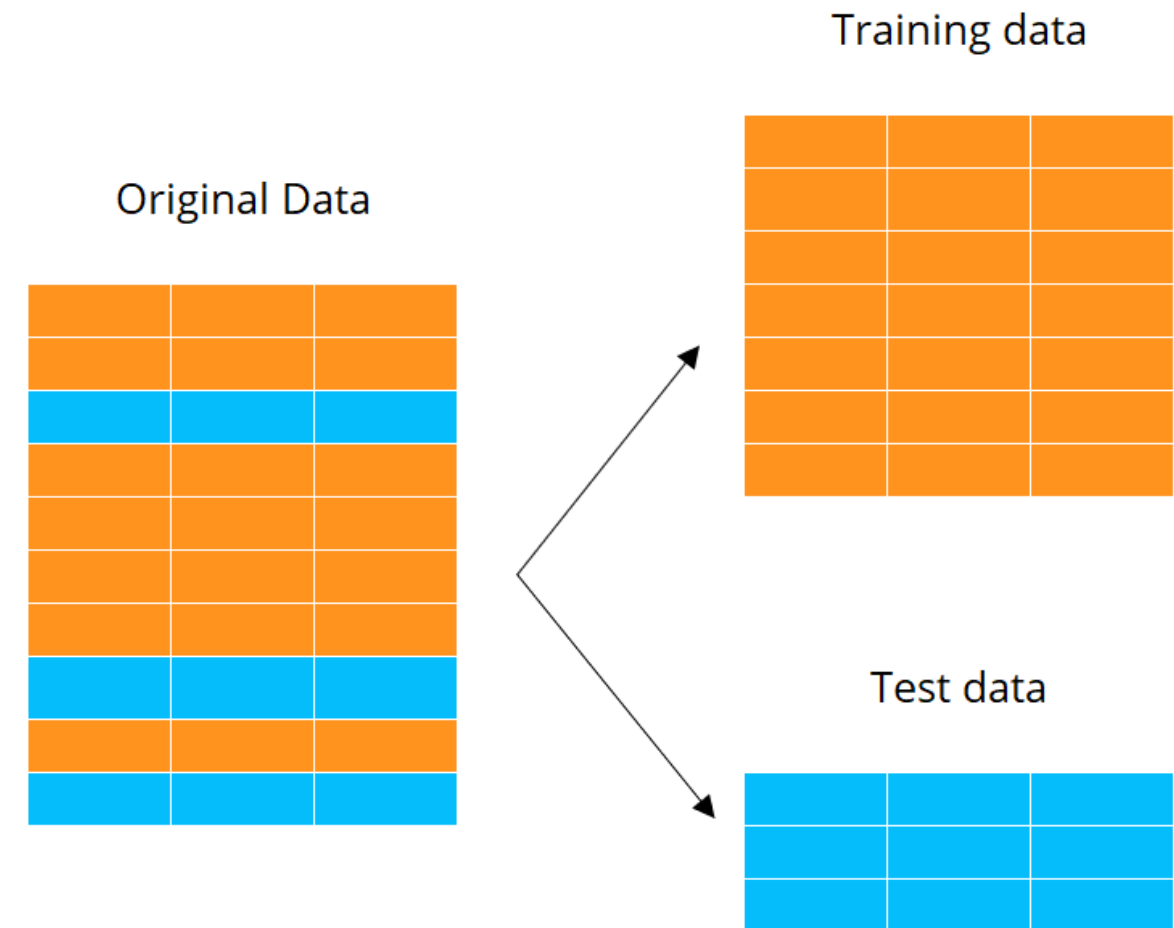
Training and test datasets

Creating training and test datasets is the first step in the modeling process

- Guards against **overfitting**
 - Training data is used for model fitting
 - Test data is used for model evaluation

Downside


- Only **one** estimate of model performance



K-fold cross validation

Resampling technique for exploring model performance

- Provides K estimates of model performance during the model fitting process

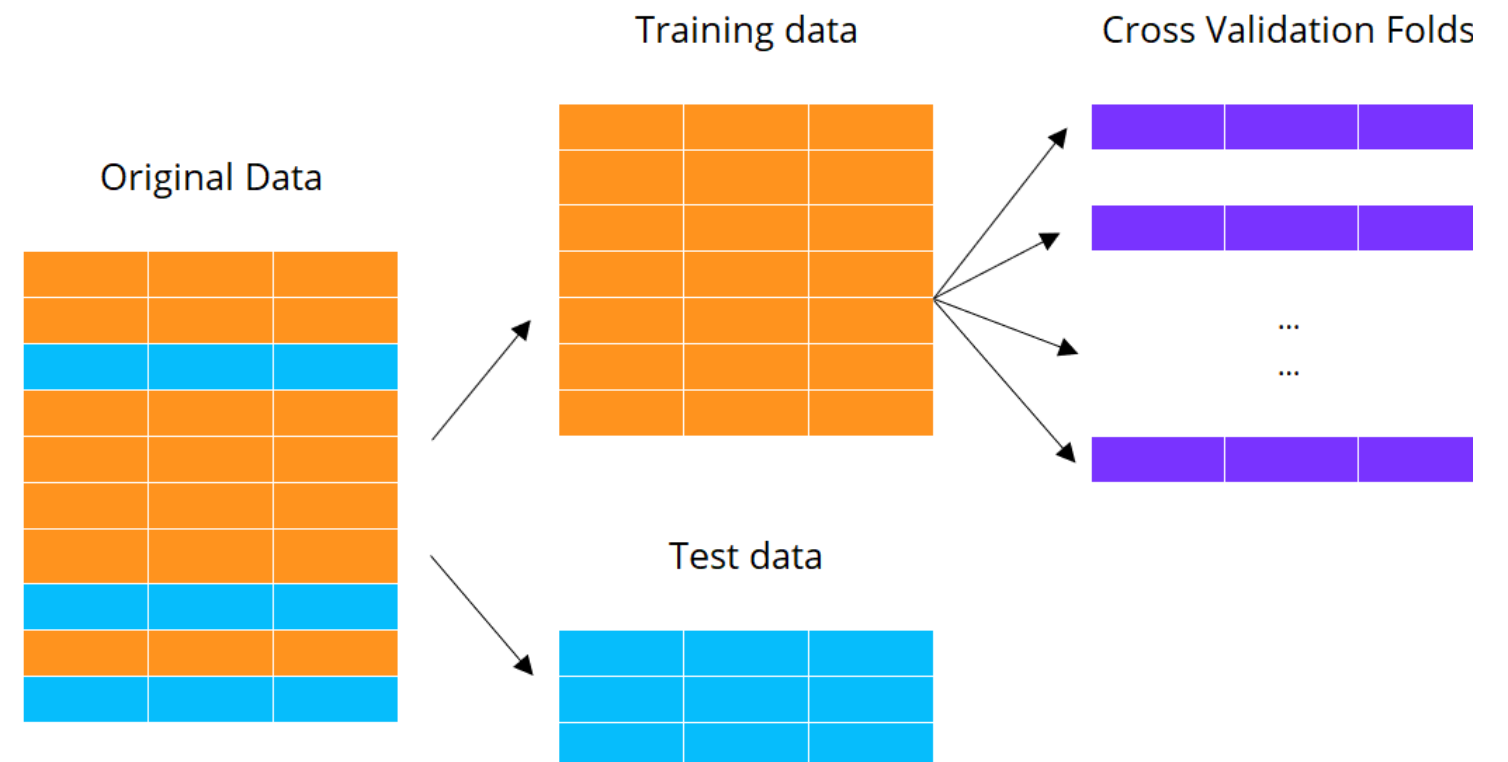


K-fold cross validation

Resampling technique for exploring model performance

often comparing different models (e.g., logistic regression & decision tree)

- Provides K estimates of model performance during the model fitting process
- Training data is randomly partitioned into K sets of roughly equal size
- Folds are used to perform K iterations of model fitting and evaluation




Machine learning with cross validation

Performing 5-fold cross validation

- Five iterations of model training and evaluation

Cross Validation Folds




Machine learning with cross validation

Performing 5-fold cross validation

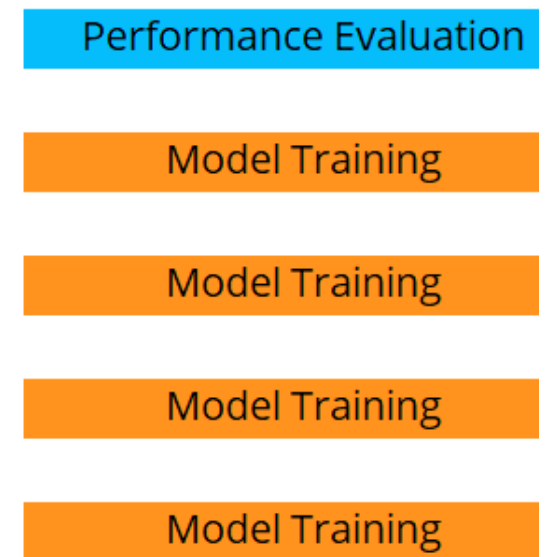
- Five iterations of model training and evaluation
- Iteration 1
 - Fold 1 reserved for model evaluation and folds 2 through 5 for model training

Cross Validation Folds



Modeling
→

Iteration 1




Machine learning with cross validation

Performing 5-fold cross validation

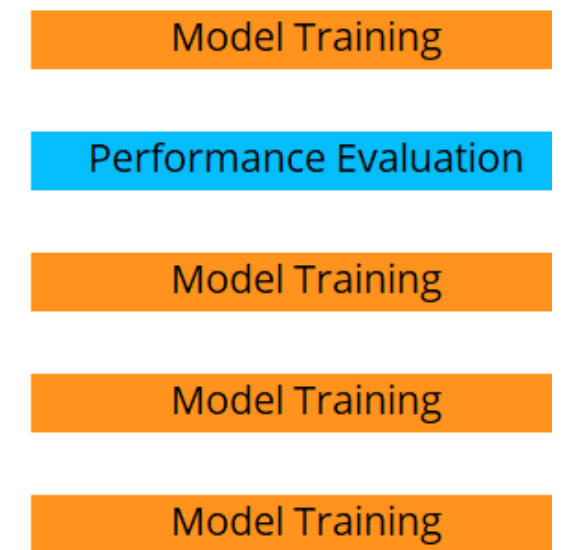
- Five iterations of model training and evaluation
- Iteration 1
 - Fold 1 reserved for model evaluation and folds 2 through 5 for model training
- Iteration 2
 - Fold 2 reserved for model evaluation

Cross Validation Folds



Modeling
→

Iteration 2




Machine learning with cross validation

Performing 5-fold cross validation

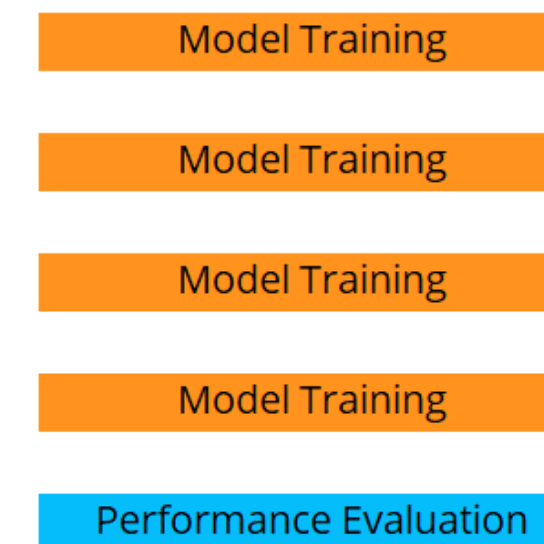
- Five iterations of model training and evaluation
- Iteration 1
 - Fold 1 reserved for model evaluation and folds 2 through 5 for model training
- Iteration 2
 - Fold 2 reserved for model evaluation

Cross Validation Folds



Modeling
→

Iteration 5



Five estimates of model performance in total

Creating cross validation folds

The `vfold_cv()` function

- Training data
- Number of folds, `v`
- Stratification variable, `strata`
- Execute `set.seed()` before `vfold_cv()` for reproducibility
- `splits`
 - **List column** with data split objects for creating fold

```
set.seed(214)
leads_folds <- vfold_cv(leads_training,
                        v = 10,
                        strata = purchased)

leads_folds
```

```
# 10-fold cross-validation using stratification
# A tibble: 10 x 2
  splits          id
  <list>         <chr>
1 <split [896/100]> Fold01
2 <split [896/100]> Fold02
3 <split [896/100]> Fold03
. ....
9 <split [897/99]>  Fold09
10 <split [897/99]> Fold10
```


Model training with cross validation

The `fit_resamples()` function

- Train a `parsnip` model or `workflow` object
- Provide cross validation folds, `resamples`
- Optional custom metric function, `metrics`
 - Default is accuracy and ROC AUC

Each metric is estimated 10 times

- One estimate per fold
- Average value in `mean` column

includes models and data pre-processing

```
leads_rs_fit <- leads_wkfl %>%  
  fit_resamples(resamples = leads_folds,  
               metrics = leads_metrics)  
  
leads_rs_fit %>%  
  collect_metrics()
```

```
# A tibble: 3 x 5  
  .metric .estimator mean      n std_err  
  <chr>   <chr>     <dbl> <int>   <dbl>  
1 roc_auc binary    0.823    10  0.0147  
2 sens   binary    0.786    10  0.0203  
3 spec   binary    0.855    10  0.0159
```

Detailed cross validation results

The `collect_metrics()` function

- Passing `summarize = FALSE` will provide all metric estimates for every cross validation fold
- 30 total combinations (3 metrics x 10 folds)
 - `.metric` column identifies metric
 - `.estimate` column gives estimated value for each fold

```
rs_metrics <- leads_rs_fit %>%  
  collect_metrics(summarize = FALSE)
```

```
rs_metrics
```

```
# A tibble: 30 x 4  
  id      .metric .estimator .estimate  
  <chr>   <chr>    <chr>      <dbl>  
1 Fold01 sens     binary     0.861  
2 Fold01 spec     binary     0.891  
3 Fold01 roc_auc  binary     0.885  
4 Fold02 sens     binary     0.778  
5 Fold02 spec     binary     0.969  
6 Fold02 roc_auc  binary     0.885  
# ... with 24 more rows
```

Summarizing cross validation results

The `collect_metrics()` function returns a tibble

- Results can be summarized with `dplyr`
 - Start with `rs_metrics`
 - Form groups by `.metric` values
 - Calculate summary statistics with `summarize()`

```
rs_metrics %>%  
  group_by(.metric) %>%  
  summarize(min = min(.estimate),  
            median = median(.estimate),  
            max = max(.estimate),  
            mean = mean(.estimate),  
            sd = sd(.estimate))
```

```
# A tibble: 3 x 6  
  .metric    min median    max  mean    sd  
  <chr>    <dbl> <dbl> <dbl> <dbl> <dbl>  
1 roc_auc 0.758  0.806 0.885 0.823 0.0466  
2 sens    0.667  0.792 0.861 0.786 0.0642  
3 spec    0.810  0.843 0.969 0.855 0.0502
```

Cross validation methodology

Models trained with `fit_resamples()` **are not** able to provide predictions on new data sources

- `predict()` function does not accept resample objects

Purpose of `fit_resample()`

- Explore and compare the performance profile of different model types
- Select best performing model type and focus on model fitting efforts

```
predict(leads_rs_fit,  
        new_data = leads_test)
```

```
Error in UseMethod("predict") :  
  no applicable method for 'predict' applied to  
  an object of class  
  "c('resample_results',  
    'tune_results',  
    'tbl_df',  
    'tbl', 'data.frame')"
```

Let's cross validate!

MODELING WITH TIDYMODELS IN R

Hyperparameter tuning

MODELING WITH TIDYMODELS IN R



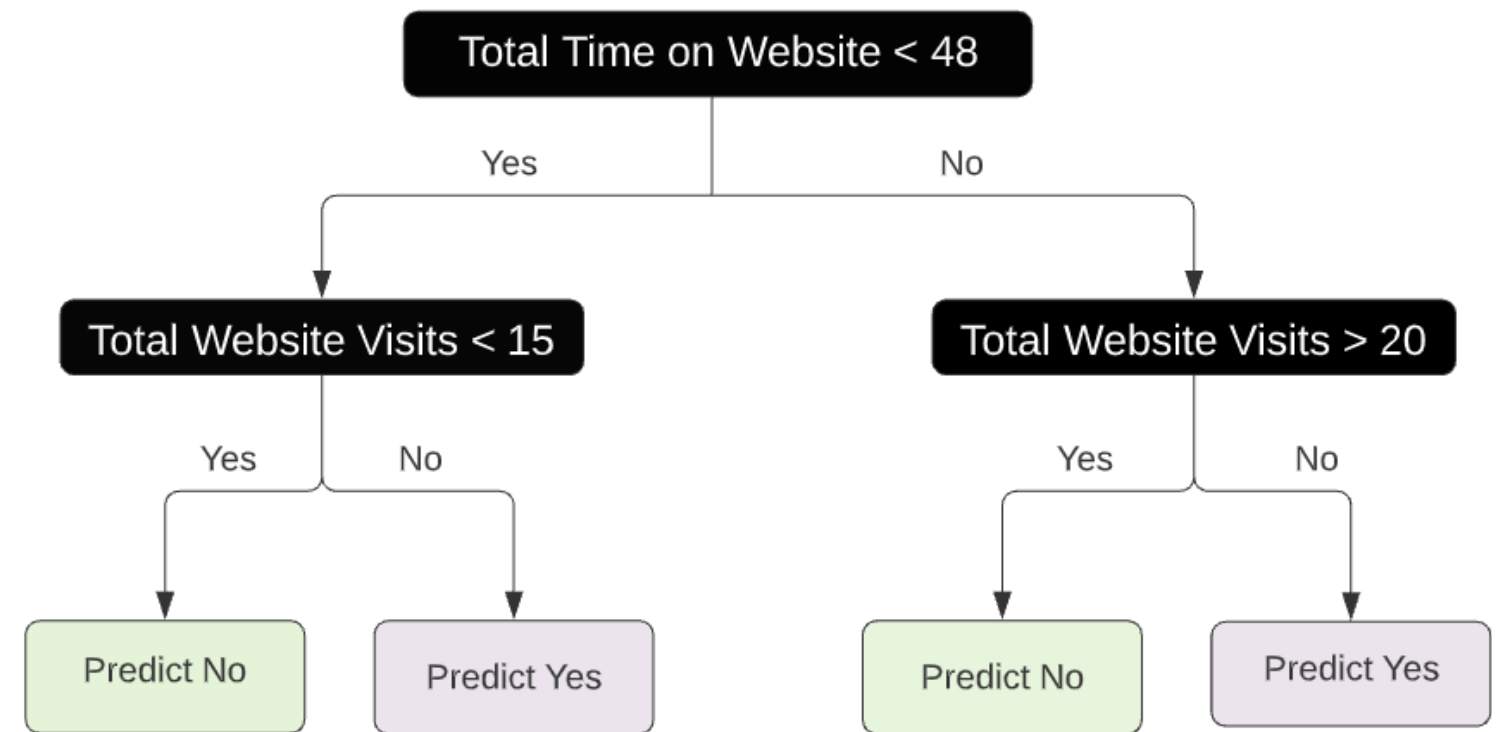
David Svancer
Data Scientist

Hyperparameters

Model parameters whose values are set prior to model training and control model complexity

`parsnip` decision tree

- `cost_complexity`
 - Penalizes large number of terminal nodes
- `tree_depth`
 - Longest path from root to terminal node
- `min_n`
 - Minimum data points required in a node for further splitting



Default hyperparameter values

`decision_tree()` function sets default hyperparameter values

- `cost_complexity` is set to 0.01
- `tree_depth` is set to 30
- `min_n` is set to 20

These may not be the best values for all datasets

- **Hyperparameter tuning**
 - Process of using cross validation to find the optimal set of hyperparameter values

```
dt_model <- decision_tree() %>%  
  set_engine('rpart') %>%  
  set_mode('classification')
```


Labeling hyperparameters for tuning

The `tune()` function from the `tune` package

- To label hyperparameters for tuning, set them equal to `tune()` in `parsnip` model specification
- Creates model object with tuning parameters
 - Will let other functions know that they need to be optimized

```
dt_tune_model <- decision_tree(cost_complexity = tune(),  
                               tree_depth = tune(),  
                               min_n = tune()) %>%  
  
  set_engine('rpart') %>%  
  set_mode('classification')  
  
dt_tune_model
```

Decision Tree Model Specification (classification)

Main Arguments:

```
cost_complexity = tune()  
tree_depth = tune()  
min_n = tune()
```

Computational engine: rpart

Creating a tuning workflow

`workflow` objects can be easily updated

- Prior `leads_wkfl`
 - Feature engineering steps for lead scoring data and decision tree model with default hyperparameters
- Pass `leads_wkfl` to `update_model()` and provide new decision tree model with tuning parameters

```
leads_tune_wkfl <- leads_wkfl %>%  
  update_model(dt_tune_model)
```

```
leads_tune_wkfl
```

```
== Workflow =====  
Preprocessor: Recipe  
Model: decision_tree()  
-- Preprocessor -----  
3 Recipe Steps  
* step_corr()  
* step_normalize()  
* step_dummy()  
-- Model -----  
Decision Tree Model Specification (classification)  
Main Arguments: cost_complexity = tune()  
                  tree_depth = tune()  
                  min_n = tune()  
Computational engine: rpart
```

Grid search

Most common method for tuning hyperparameters

- Generate a grid of unique combinations of hyperparameter values
 - For each combination, use cross validation to estimate model performance
- Choose best performing combination

cost_complexity	tree_depth	min_n
0.001	20	35
0.001	20	15
0.001	35	35
0.001	35	15
0.2	20	35
...

Identifying hyperparameters

The `parameters()` function from the `dials` package

- Takes a `parsnip` model object
- Returns a tibble with the hyperparameters labeled by the `tune()` function, if any
 - Used for generating tuning grids with the `dials` package

```
parameters(dt_tune_model)
```

Collection of 3 parameters for tuning

identifier	type	object
cost_complexity	cost_complexity	nparam[+]
tree_depth	tree_depth	nparam[+]
min_n	min_n	nparam[+]

Random grid

Generating random combinations

- This method tends to provide greater chances of finding optimal hyperparameter values

The `grid_random()` function

- First argument is the results of the `parameters()` function
- `size` sets the number of random combinations to generate
 - Execute `set.seed()` function before `grid_random()` for reproducibility

```
set.seed(214)
grid_random(parameters(dt_tune_model),
             size = 5)
```

```
# A tibble: 5 x 3
  cost_complexity tree_depth min_n
          <dbl>         <int> <int>
1  0.00000000758          14     39
2  0.0243              5      34
3  0.000000443          11      8
4  0.0000000600           3      5
5  0.00380              5     36
```

Saving a tuning grid

First step in hyperparameter tuning

- Create and save a tuning grid
- `dt_grid` contains 5 random combinations of hyperparameter values

```
set.seed(214)
dt_grid <- grid_random(parameters(dt_tune_model),
                        size = 5)
```

```
dt_grid
```

```
# A tibble: 5 x 3
  cost_complexity tree_depth min_n
      <dbl>         <int> <int>
1 0.0000000758         14     39
2 0.0243             5     34
3 0.00000443          11      8
4 0.000000600          3      5
5 0.00380             5     36
```

Hyperparameter tuning with cross validation

The `tune_grid()` function performs hyperparameter tuning

Takes the following arguments:

- `workflow` or `parsnip` model
- Cross validation object, `resamples`
- Tuning grid, `grid`
- Optional `metrics` function

Returns tibble of results

- `.metrics`
 - List column with results for each fold

```
dt_tuning <- leads_tune_wkfl %>%  
  tune_grid(resamples = leads_folds,  
            grid = dt_grid,  
            metrics = leads_metrics)
```

dt_tuning

```
# Tuning results  
# 10-fold cross-validation using stratification  
# A tibble: 10 x 4  
  splits          id    .metrics      ..  
  <list>         <chr>  <list>      ..  
<split [896/100]> Fold01 <tibble [15 x 7]> ..  
.....  
<split [897/99]>  Fold09 <tibble [15 x 7]> ..  
<split [897/99]>  Fold10 <tibble [15 x 7]> ..
```

Exploring tuning results

The `collect_metrics()` function provides summarized results by default

- Average estimated metric values across all folds per combination

```
dt_tuning %>%  
  collect_metrics()
```

```
# A tibble: 15 x 9  
  cost_complexity tree_depth min_n .metric .estimator mean      n std_err .config  
          <dbl>      <int> <int> <chr>   <chr>    <dbl> <int>   <dbl> <chr>  
1  0.0000000758      14     39 roc_auc binary  0.827    10 0.0147 Model1  
2  0.0000000758      14     39 sens    binary  0.728    10 0.0277 Model1  
3  0.0000000758      14     39 spec    binary  0.865    10 0.0156 Model1  
4  0.0243           5     34 roc_auc binary  0.823    10 0.0147 Model2  
·      ······      ··     ··  ······  ······  ······  ··  ······  ······  
14 0.00380          5     36 sens    binary  0.747    10 0.0209 Model5  
15 0.00380          5     36 spec    binary  0.858    10 0.0161 Model5
```


Let's get tuning!

MODELING WITH TIDYMODELS IN R

Selecting the best model

MODELING WITH TIDYMODELS IN R



David Svancer
Data Scientist

Detailed tuning results

The `collect_metrics()` function provides summarized results by default

- Passing `summarize = FALSE` will provide all hyperparameter tuning results

```
dt_tuning %>%  
  collect_metrics(summarize = FALSE)
```

```
# A tibble: 150 x 8  
  id      cost_complexity tree_depth min_n .metric ... .estimate .config  
<chr>      <dbl>          <int> <int> <chr> ...   <dbl>    <chr>  
Fold01  0.00000000758         14      39  sens ...    0.75    Model1  
Fold01  0.00000000758         14      39  spec ...    0.906    Model1  
Fold01  0.00000000758         14      39 roc_auc ...    0.888    Model1  
.....  
Fold10  0.00380              5      36 roc_auc ...    0.789    Model5
```

Exploring tuning results

Selecting `summarise = FALSE` within `collect_metrics()` returns a tibble

- Easy to explore results with `dplyr`
- Exploring ROC AUC
 - Select `roc_auc` metric
 - Form groups by `id` column
 - Calculate `.estimate` summary statistics

```
dt_tuning %>%  
  collect_metrics(summarize = FALSE) %>%  
  filter(.metric == 'roc_auc') %>%  
  group_by(id) %>%  
  summarize(min_roc_auc = min(.estimate),  
            median_roc_auc = median(.estimate),  
            max_roc_auc = max(.estimate))
```

```
# A tibble: 10 x 4  
  id      min_roc_auc median_roc_auc max_roc_auc  
<chr>    <dbl>         <dbl>         <dbl>  
Fold01  0.830         0.885         0.888  
Fold02  0.857         0.882         0.885  
Fold03  0.818         0.836         0.836  
.....  
Fold10  0.762         0.790         0.813
```

Viewing the best performing models

The `show_best()` function

- Displays the top `n` performing models based on average value of `metric`
- `Model1` is the winner

```
dt_tuning %>%  
  show_best(metric = 'roc_auc', n = 5)
```

```
# A tibble: 5 x 9  
cost_complexity tree_depth min_n .metric .estimator mean n std_err .config  
      <dbl>         <int> <int>   <chr>   <chr>    <dbl> <int> <dbl>   <chr>  
0.0000000758      14      39  roc_auc binary    0.827   10  0.0147 Model1  
0.00380           5      36  roc_auc binary    0.825   10  0.0146 Model5  
0.0243            5      34  roc_auc binary    0.823   10  0.0147 Model2  
0.00000443        11       8  roc_auc binary    0.816   10  0.00786 Model3  
0.000000600        3       5  roc_auc binary    0.814   10  0.0131 Model4
```

Selecting a model

The `select_best()` function

- Pass `dt_tuning` results to `select_best()`
- Select the `metric` on which to evaluate performance

Returns a tibble with the best performing model and hyperparameter values

```
best_dt_model <- dt_tuning %>%  
  select_best(metric = 'roc_auc')
```

```
best_dt_model
```

```
# A tibble: 1 x 4  
cost_complexity tree_depth min_n .config  
      <dbl>         <int> <int>   <chr>  
1 0.00000000758         14    39 Model1
```

Finalizing the workflow

The `finalize_workflow()` function will finalize a `workflow` that contains a model object with tuning parameters

- Pass `workflow` object
- A tibble with one row of final model hyperparameter values
 - Column names must match hyperparameters in model object

Returns a `workflow` object with set hyperparameter values

```
final_leads_wkfl <- leads_tune_wkfl %>%  
  finalize_workflow(best_dt_model)  
final_leads_wkfl
```

```
== Workflow =====  
Preprocessor: Recipe  
Model: decision_tree()  
-- Preprocessor -----  
3 Recipe Steps  
* step_corr()  
* step_normalize()  
* step_dummy()  
-- Model -----  
Decision Tree Model Specification (classification)  
Main Arguments:  
  cost_complexity = 0.0000000758  
  tree_depth = 14  
  min_n = 39  
Computational engine: rpart
```

Model fitting

Finalized `workflow` object can be trained with `last_fit()` and original data split object, `leads_split`

```
leads_final_fit <- final_leads_wkfl %>%  
  last_fit(split = leads_split)  
  
leads_final_fit %>%  
  collect_metrics()
```

Behind the scenes

- Training and test datasets created
- `recipe` trained and applied
- **Tuned decision tree** trained with entire training dataset
- Predictions and metrics on test data

```
# A tibble: 2 x 3  
  .metric .estimator .estimate  
  <chr>    <chr>         <dbl>  
1 accuracy binary        0.771  
2 roc_auc  binary        0.793
```


Let's practice!

MODELING WITH TIDYMODELS IN R

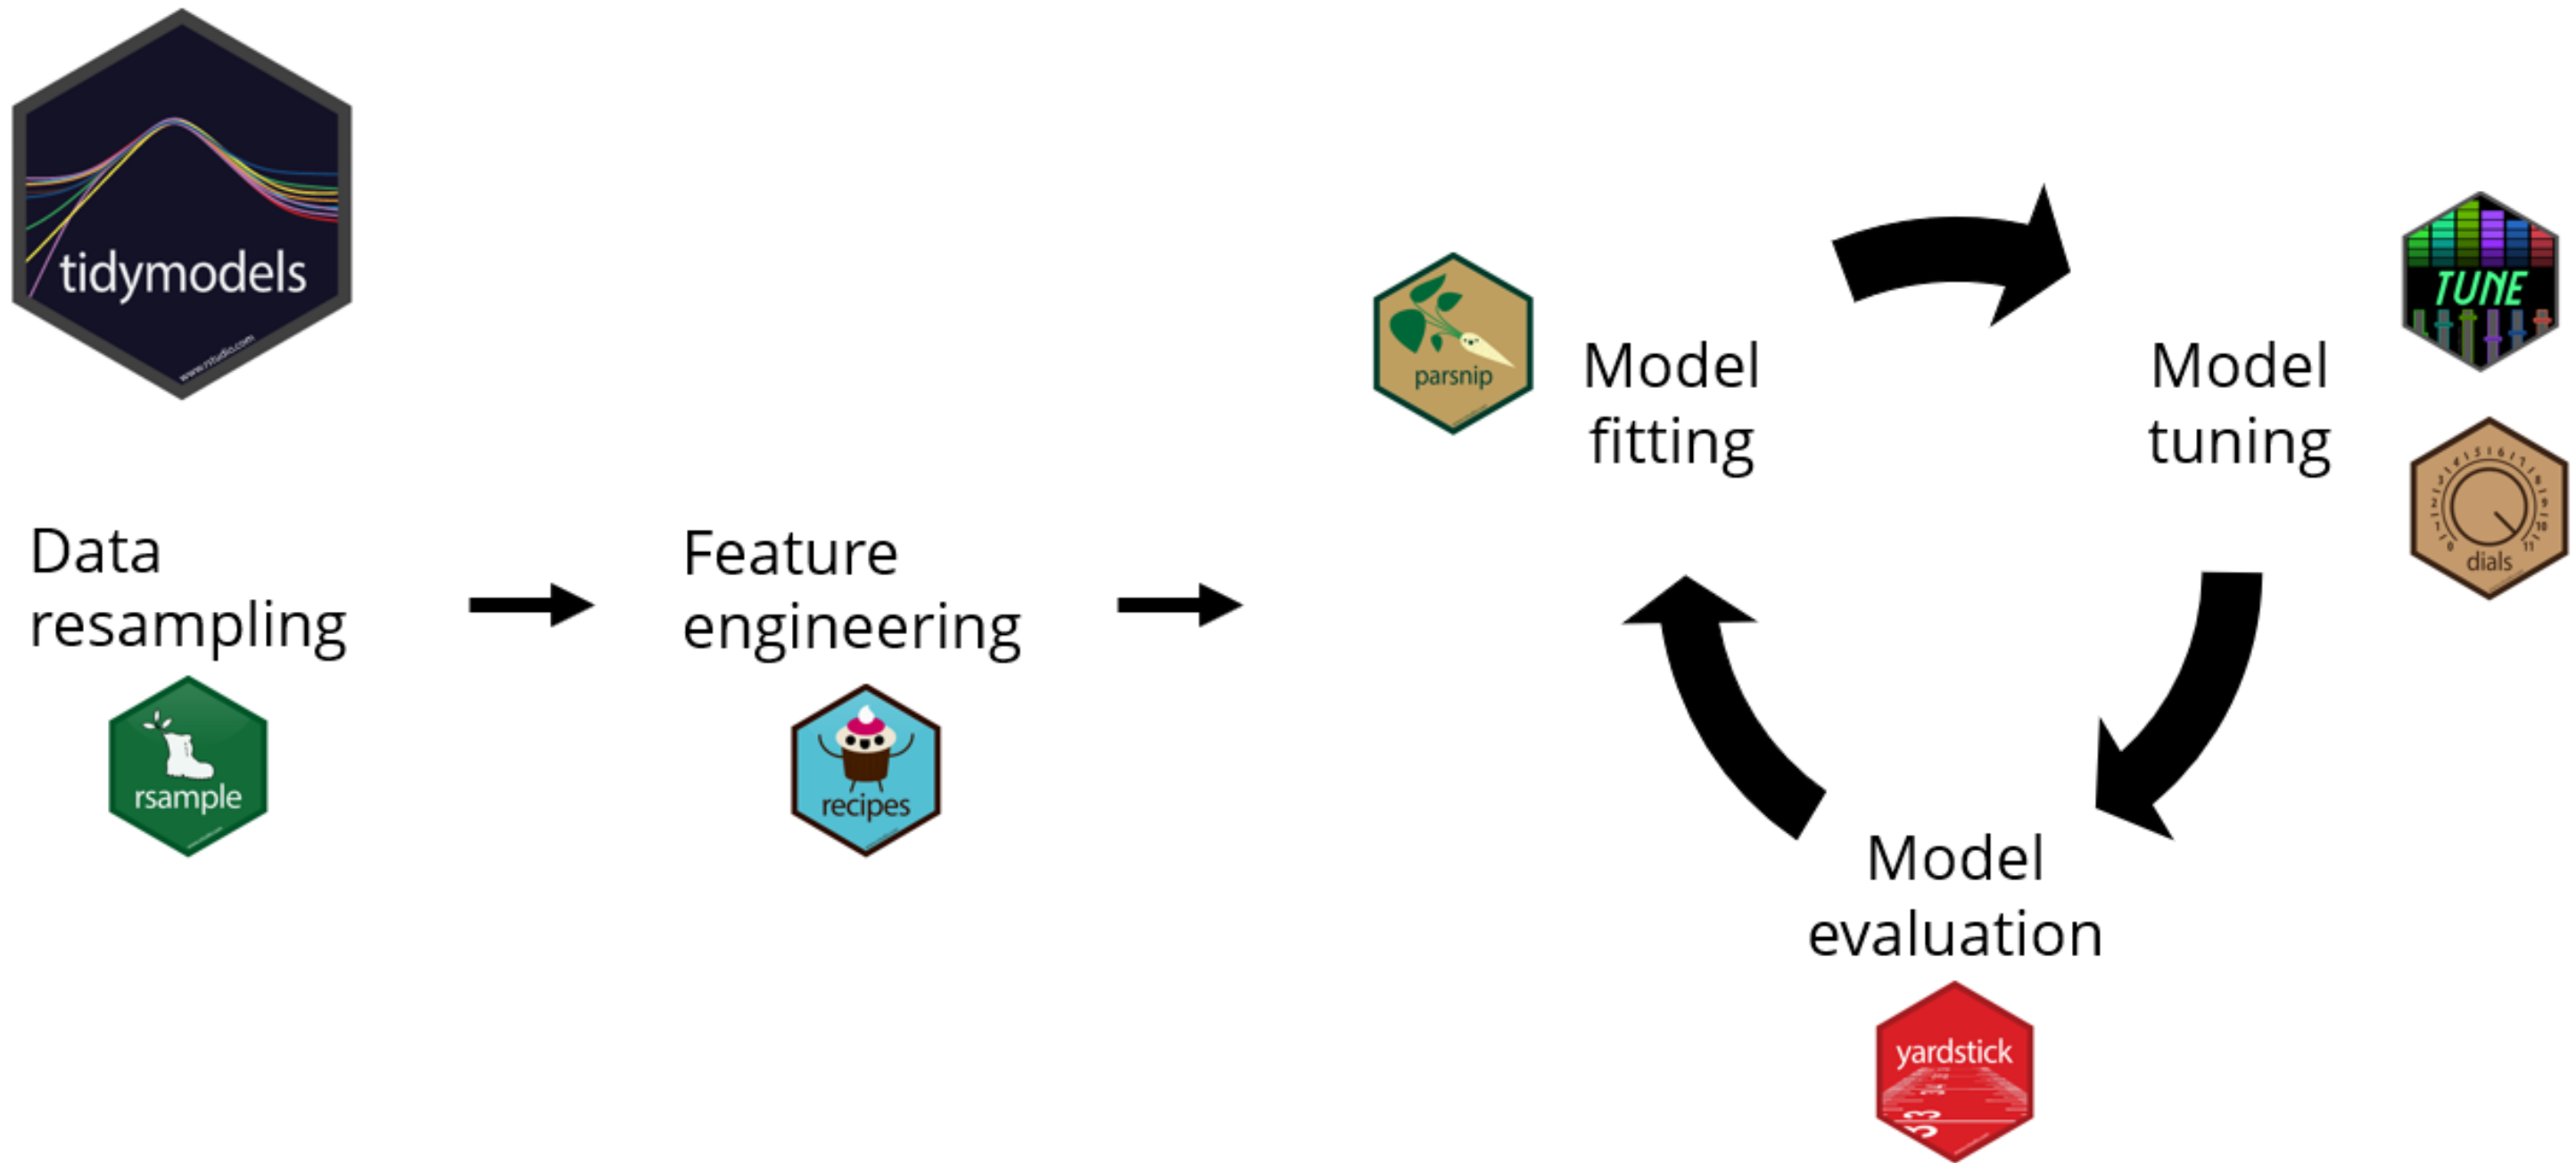
Congratulations!

MODELING WITH TIDYMODELS IN R



David Svancer
Data Scientist

The tidymodels ecosystem




Regression modeling

Specifying models with `parsnip`

Training and evaluating linear regression models





R-Squared Plot



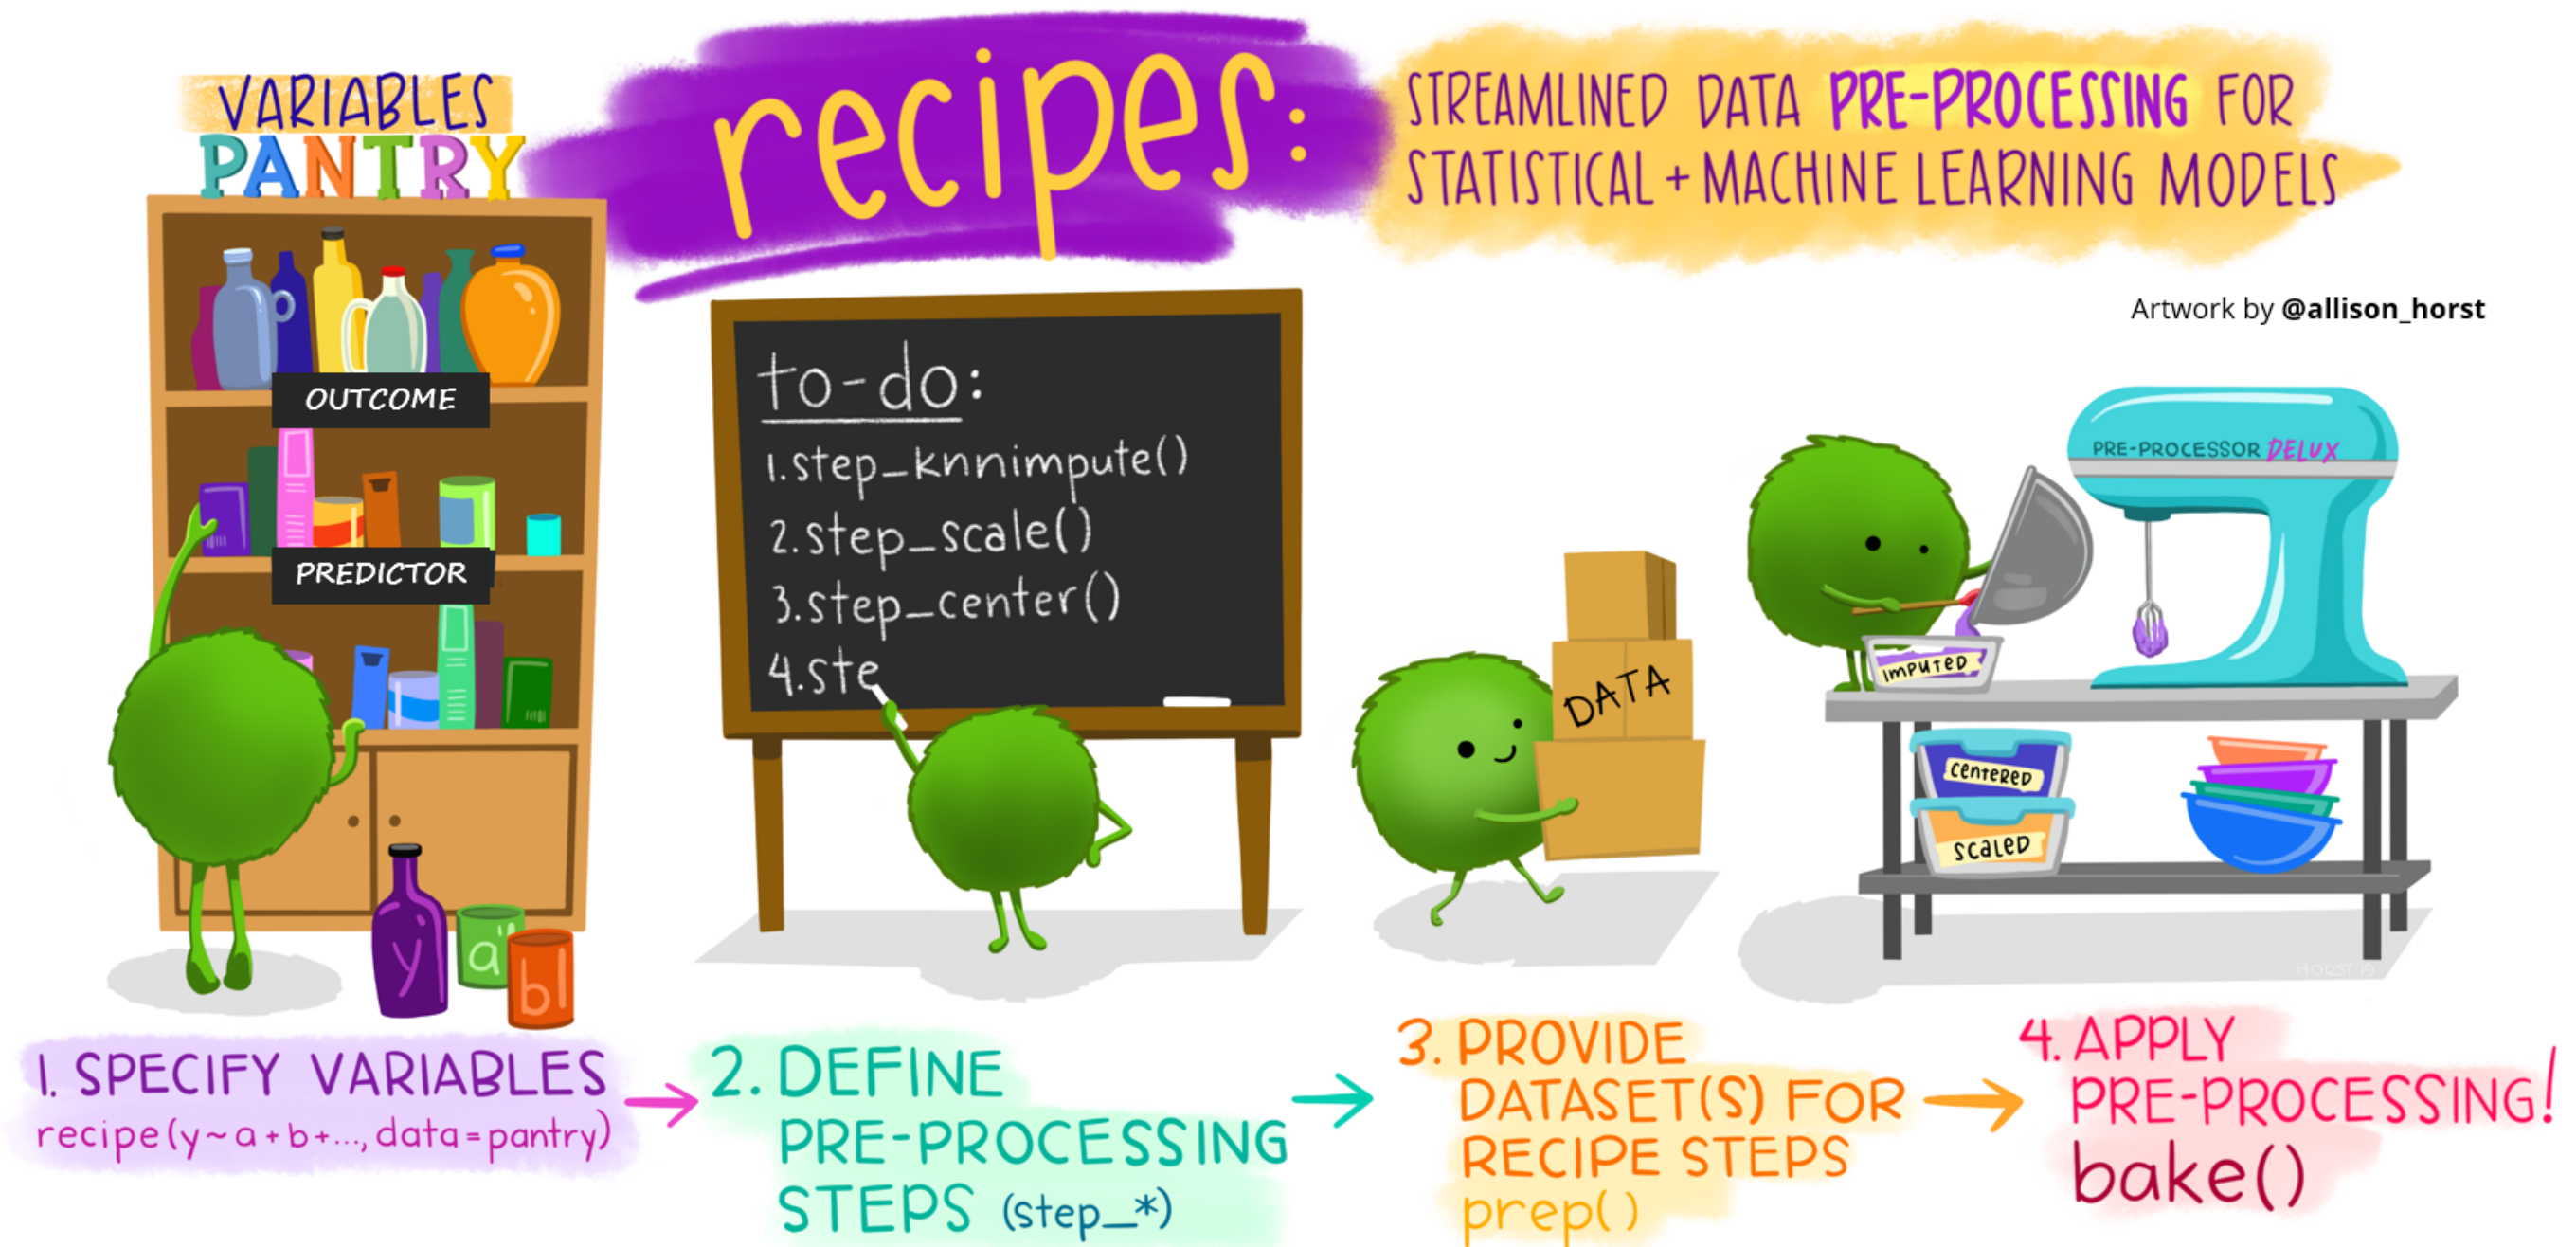
Classification modeling

Logistic regression with `logistic_reg()`

Evaluating classification performance with confusion matrices and ROC curves



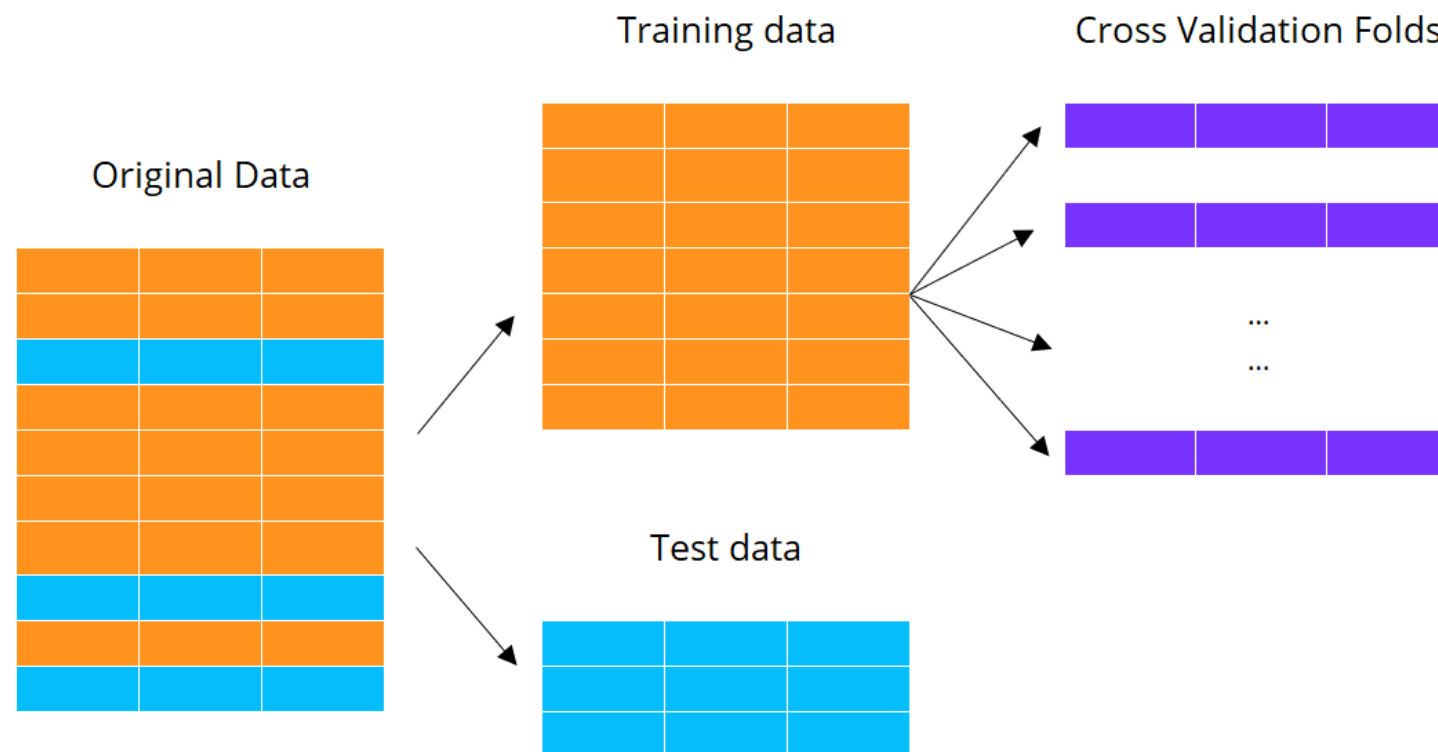
Feature engineering



Fine tuning models with cross validation

Model performance profiles with cross validation and `fit_resamples()`

- Hyperparameter tuning with grid search
- Finalizing model workflows



cost_complexity	tree_depth	min_n
0.001	20	35
0.001	20	15
0.001	35	35
0.001	35	15
0.2	20	35
...

Thank you!

MODELING WITH TIDYMODELS IN R